

# Contents

- [Import packages and classes](#)
- [Load dataset](#)
- [Select variables & Drop NaN values](#)
- [Train-test split](#)
- [Sklean Linear Regression](#)
- [Prediction](#)
- [Evaluation](#)

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn import metrics
import statsmodels.api as sm
```

```
data_raw = pd.read_csv('../datasets/Greece - Agriculture and Rural
Development/agriculture-and-rural-development_grc.csv')
pd.set_option('display.max_columns', 4)
data_raw.head()
```

	Agricultural land (% of land area)	Agricultural land (sq. km)	...	Rural population growth (annual %)	Surface area (sq. km)
0	NaN	NaN	...	NaN	NaN
1	69.123351	89100.0	...	-0.387316	131960.0
2	69.061288	89020.0	...	-1.462143	131960.0
3	69.984484	90210.0	...	-1.718278	131960.0
4	69.751746	89910.0	...	-1.758920	131960.0

5 rows × 23 columns

```
variable_x = "Crop production index (2014-2016 = 100)"
variable_y = "Cereal production (metric tons)"
data = data_raw[[variable_x, variable_y]].dropna().reset_index(drop=True)
data.head()
```

	Crop production index (2014-2016 = 100)	Cereal production (metric tons)
0	62.570000	2243876.0
1	50.700001	2426843.0
2	57.430000	2122537.0
3	57.150002	2874641.0
4	61.770000	2940922.0

```
data.shape
```

```
(58, 2)
```

```
X_train, X_test, y_train, y_test = train_test_split(data[variable_x].to_numpy(),
data[variable_y].to_numpy(), test_size=0.2, random_state=42)
```

```
model = LinearRegression().fit(X_train.reshape(-1, 1), y_train)
```


```
print(f"Intercept: {model.intercept_}")  
print(f"Slope: {model.coef_}")
```

```
Intercept: -338797.3306638906  
Slope: [51647.36492435]
```

```
y_pred = np.around(model.predict(X_test.reshape(-1, 1)), 1)  
y_pred
```

```
array([2892778.3, 2960952.8, 5072813.5, 3795574.4, 5177141.4, 4196357.7,  
       4998957.8, 4700952.6, 4551691.8, 3424746.3, 4813543.9, 2612849.7])
```

```
fig, ax = plt.subplots()  
plt.scatter(X_train, y_train, color='red')  
ax.axline((X_test[1], y_pred[1]), slope=model.coef_[0])  
  
plt.title('Linear regression - Cereal production', fontsize=16)  
plt.xlabel(variable_x, fontsize=14)  
plt.ylabel(variable_y, fontsize=14)  
  
plt.tight_layout()  
plt.show()
```



```
print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))  
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))  
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,  
y_pred)))
```

```
Mean Absolute Error: 300467.2  
Mean Squared Error: 142457458025.50165  
Root Mean Squared Error: 377435.36933560116
```

```
print("R-Square value:", metrics.r2_score(y_test, y_pred))
```

```
R-Square value: 0.8563623217388497
```

```
X_stat = sm.add_constant(X_train)  
regsummary = sm.OLS(y_train, X_stat).fit()  
regsummary.summary()
```

<b>Dep. Variable:</b>	y	<b>R-squared:</b>	0.744
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.738
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	127.6
<b>Date:</b>	Fri, 16 Dec 2022	<b>Prob (F-statistic):</b>	1.36e-14
<b>Time:</b>	14:35:49	<b>Log-Likelihood:</b>	-668.16
<b>No. Observations:</b>	46	<b>AIC:</b>	1340.
<b>Df Residuals:</b>	44	<b>BIC:</b>	1344.
<b>Df Model:</b>	1		
<b>Covariance Type:</b>	nonrobust		

OLS Regression Results

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	-3.388e+05	4.26e+05	-0.796	0.430	-1.2e+06	5.19e+05
<b>x1</b>	5.165e+04	4571.889	11.297	0.000	4.24e+04	6.09e+04
<b>Omnibus:</b>	0.370	<b>Durbin-Watson:</b>	1.986			
<b>Prob(Omnibus):</b>	0.831	<b>Jarque-Bera (JB):</b>	0.538			
<b>Skew:</b>	-0.082	<b>Prob(JB):</b>	0.764			
<b>Kurtosis:</b>	2.496	<b>Cond. No.</b>	534.			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
fig, ax = plt.subplots()
x_axis = np.arange(12)

plt.bar(x_axis-0.2, y_test, width=0.4, label = "Actual")
plt.bar(x_axis+0.2, y_pred, width=0.4, label = "Predicted")

plt.xlabel("Sample", fontsize=14)
plt.ylabel("Count", fontsize=14)

plt.legend()

plt.tight_layout()
plt.show()
```

